



Deep Learning


Lecture 3



Gradient Descent


- It is an optimization algorithm used in machine learning and deep learning to minimize a cost function by iteratively updating model parameters (weights and biases)
- The goal is to find the best parameters that reduce prediction errors.

Gradient Descent

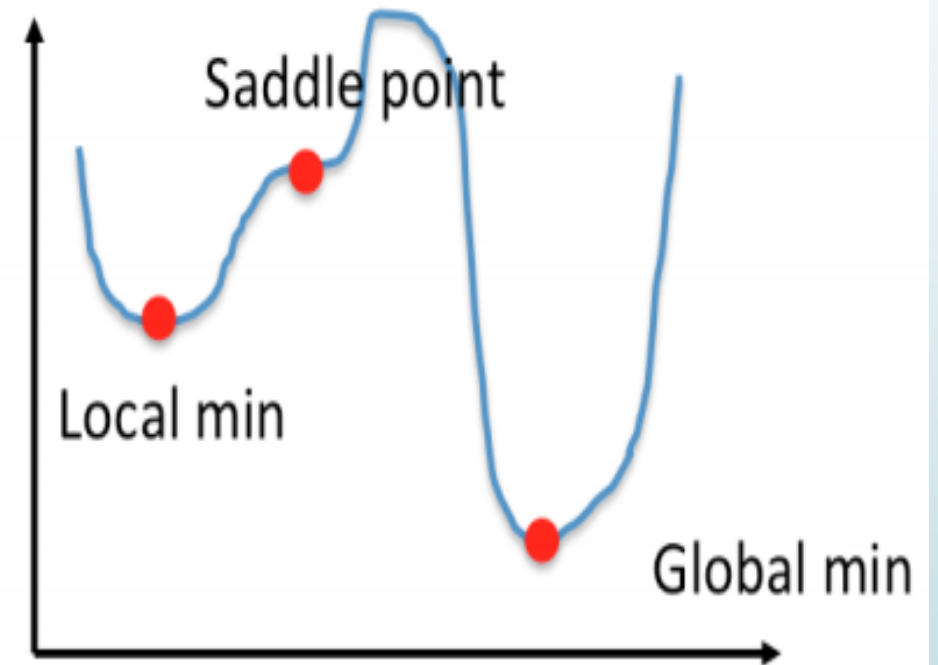


Convex VS Non-Convex

Convex



Non-Convex



<https://ai.plainenglish.io/navigating-the-terrain-convex-vs-non-convex-functions-in-optimization-86812e9a1989>



How It Works (Intuition)

Say you're at the top of a mountain and want to reach the lowest point (global minimum).

- Each step you take is based on the **steepness of the slope (gradient)** at your current position.
- A large step (**high learning rate**) might cause you to jump over the minimum
- A small step (**low learning rate**) will make progress slow but steady



Components of GD

1. **Cost Function ($J(\theta)$)** – Measures the error between predicted and actual values.
2. **Gradient ($\nabla J(\theta)$)** – Direction and magnitude of change in the cost function.
3. **Learning Rate (α)** – Controls the step size for updates.

Mathematical Intuition

1. Cost Function ($J(\theta)$)

Gradient Descent optimizes a cost function $\mathbf{J}(\boldsymbol{\theta})$ to minimize the error. For example, in linear regression:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$

Mathematical Intuition

where:

- $h_{\theta}(x) = \theta_0 + \theta_1 x$ (hypothesis/prediction)
- m = number of training samples
- y_i = actual value
- $h_{\theta}(x_i)$ = predicted value

➡ Our goal is to find **θ (theta values)** that minimize **$J(\theta)$** .

Example: Finding the Minimum of a Simple Function

► We use the function:

$$J(\theta) = (\theta - 3)^2$$

The **global minimum** is at $\theta=3$.

We start with a **random initial θ value**.

We update θ using the **Gradient Descent formula**.

$$\theta = \theta - \alpha \frac{dj}{d\theta}$$

Example: Finding the Minimum of a Simple Function

Where:

$$\frac{dj}{d\theta} = 2(\theta - 3)$$

► Let's assume:

- Initial $\theta = -5$
- Learning rate $\alpha = 0.1$
- We run for 5 iterations

Example: Finding the Minimum of a Simple Function

Iteration	θ	Gradient $2(\theta - 3)$	Update $\theta - 0.1 \times \text{Gradient}$	Cost $(\theta - 3)^2$
0	-5	$2(-5 - 3) = -16$	$-5 - 0.1(-16) = -3.4$	$(-5 - 3)^2 = 64$
1	-3.4	$2(-3.4 - 3) = -12.8$	$-3.4 - 0.1(-12.8) = -2.12$	$(-3.4 - 3)^2 = 42.25$
2	-2.12	$2(-2.12 - 3) = -10.24$	$-2.12 - 0.1(-10.24) = -1.096$	$(-2.12 - 3)^2 = 25.92$
3	-1.096	$2(-1.096 - 3) = -8.192$	$-1.096 - 0.1(-8.192) = -0.2768$	$(-1.096 - 3)^2 = 16.11$
4	-0.2768	$2(-0.2768 - 3) = -6.5536$	$-0.2768 - 0.1(-6.5536) = 0.37856$	$(-0.2768 - 3)^2 = 10.65$



Observations

1. **θ moves towards 3** in each iteration.
2. The **gradient gets smaller** as θ approaches the minimum.
3. The **cost decreases** in each step.



Variants of Gradient Descent

1. **Batch Gradient Descent (BGD)** – Uses all training data at once (slow but stable)
2. **Stochastic Gradient Descent (SGD)** – Updates parameters for each data point (faster but noisy)
3. **Mini-Batch Gradient Descent** – A balance between BGD and SGD (uses small batches)



Batch Gradient Descent (BGD)

- Batch Gradient Descent (Slow but Stable)
- For example, you want to **lose weight** and reach your **ideal body weight** of **70 kg**. You **track your weight for a month**, calculate your **average weight loss**, and **then adjust** your diet and exercise. This is **slow but accurate** because you make **big adjustments based on all data at once**.



Stochastic Gradient Descent (Fast but Noisy)

- You **weigh yourself every day** and **immediately adjust** your diet based on **just that day's weight**
- This is **faster**, but sometimes **random daily fluctuations (like eating extra one day)** might cause **overreactions**



Advanced Optimization Algorithms

1. Momentum-Based GD
2. RMSProp (Root Mean Square Propagation)
3. Adam (Adaptive Moment Estimation)

Why Momentum-Based GD

In deep learning, we face non-convex optimization.



Consistent Gradient

Noisy Gradient



Momentum-Based GD

- Instead of just updating using the current gradient, we **accumulate** past gradients to add momentum.

For example, if you want to move from point A to B, you don't know where destination B is located, you ask four persons, and all of them tell you that B is located north. So, in between you increase your speed towards destination (B) by gaining confidence.

- Faster convergence avoids oscillations.
- Can overshoot if momentum is too high

Momentum-Based GD Mathematics

- Instead of just using the current gradient, Momentum GD also considers the past gradients
- This helps in smoother updates and avoids zig-zagging

$$v_t = \beta v_{t-1} + (1 - \beta) \nabla J(\theta)$$
$$\theta = \theta - \eta v_t$$

Momentum-Based GD

$$\theta = \theta - \alpha \cdot \frac{dJ}{d\theta}$$

Standard GD



Momentum-Based GD Mathematics

Where:

- v_t is the velocity (running average of gradients).
- β is a momentum term (usually 0.9).
- η is the learning rate



Example: Ball Rolling Down a Hill

Think of a **ball rolling down a hill**, it gains speed gradually instead of taking small, slow steps. This helps **faster convergence** and reduces oscillations.



RMSProp (Root Mean Square Propagation)

- ➡ RMSProp adapts the learning rate for each parameter based on the **past gradients** to avoid **oscillations**.

For example, you are **hiking on a terrain surface**. If you step too aggressively, you may fall. Instead, you take **small careful steps** where the ground is unstable (steep gradients) and **bigger steps** where the ground is flat (small gradients).





Adam (Adaptive Moment Estimation)

- Adam **combines Momentum and RMSProp**
- 

Comparisons

Optimizer	Uses Past Gradients?	Adaptive Learning Rate?	Convergence Speed
Gradient Descent	✗ No	✗ No	🚶 Slow
Momentum	✓ Yes	✗ No	🏃 Faster
RMSProp	✗ No	✓ Yes	🏃 Faster
Adam	✓ Yes	✓ Yes	🚀 Fastest





Early Stopping, and Dropout

- When training neural networks, we aim to find a model that generalizes well to unseen data
- However, deep networks often suffer from **overfitting**, where they perform well on training data but poorly on test data
- To overcome this, we use **regularization techniques** such as **Early Stopping, and Dropout**.



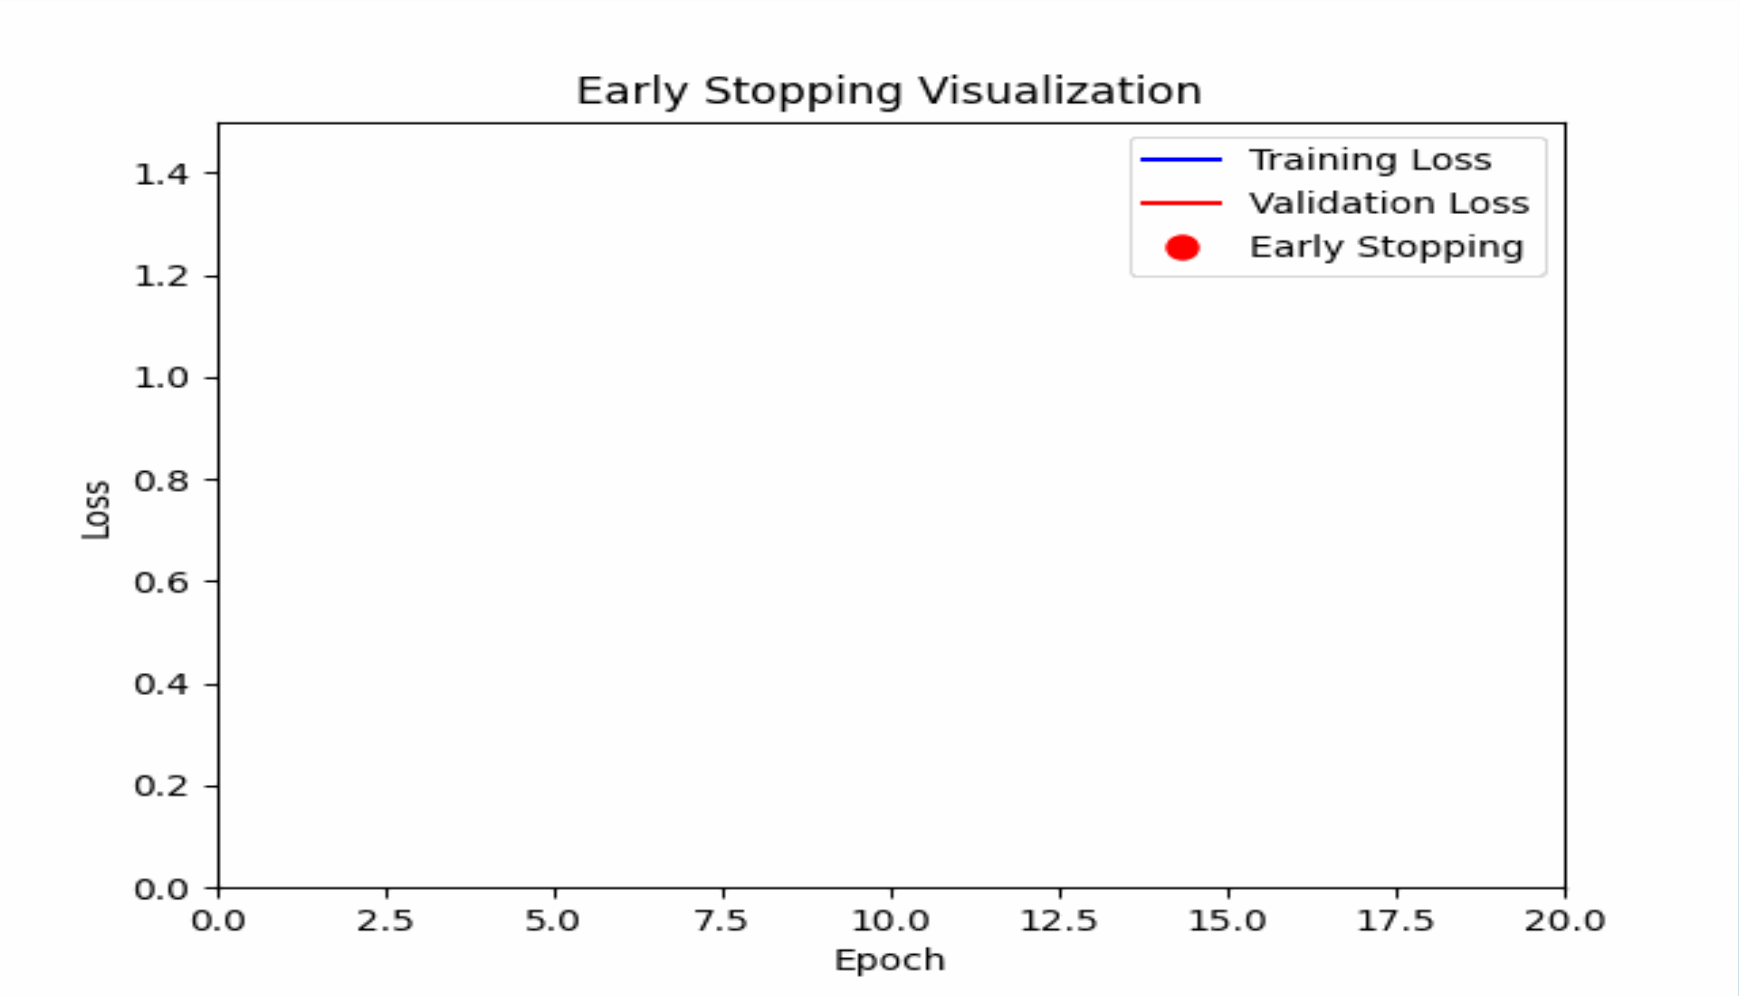
Why Early Stopping?

- As training progresses, the model starts memorizing the training data rather than learning general patterns
- If we stop too early, the model underfits; if we train too long, it overfits



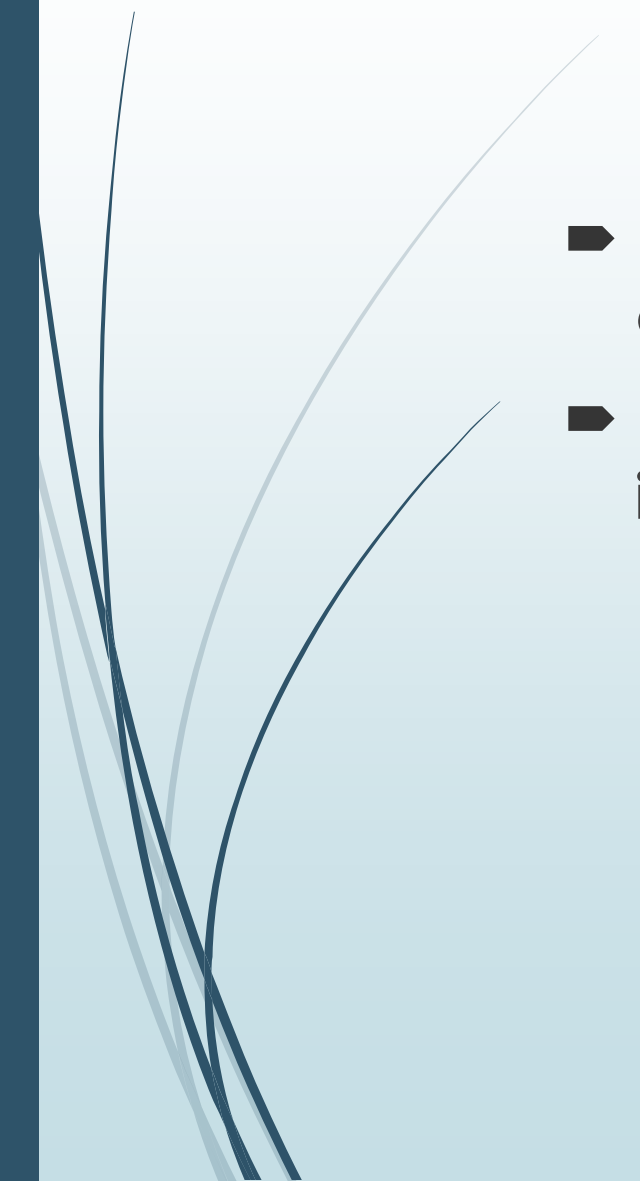
How does it work

- Monitor the model's performance on a **validation set**
- If validation loss stops improving for a defined number of epochs, stop training
- Saves computational resources and prevents overfitting.





Dropout

- In deep networks, neurons develop **dependencies** on each other, reducing their ability to generalize.
 - Dropout forces the network to **learn redundant and independent features**
- 



How does it work

- Randomly **deactivate** (set to zero) a fraction of neurons during training
 - Ensures that the network doesn't rely on specific neurons too much
- 