

Assignment-2

Name:	-----
Registration # / Roll #	-----
Subject:	-----
Class with Section and Semester #	-----
Assigned Date	20-03-2025
Due/Submission Date	27-03-2025

Due Date: Next Theory Class

Study the file annexed to this file (Annex-A: Porter Stemmer Algorithm)

Given the following paragraphs, manually apply the Porter Stemmer algorithm to reduce each word to its root form. Write down the stemmed version of the paragraphs.

Paragraph-1:

"Stemming is a technique used in natural language processing to reduce words to their base or root form. For example, the words 'running', 'runner', and 'ran' can all be reduced to the root word 'run'. Similarly, 'easily' and 'easy' share the root 'easi'. The purpose of stemming is to simplify text data for analysis by removing suffixes like '-ing', '-ed', and '-ly'."

Instructions

1. Identify each word in the paragraph.
2. Apply the rules of the Porter Stemmer algorithm manually:
 - Remove common suffixes such as -ing, -ed, -ly, etc.
 - Reduce plural forms (e.g., words → word).
 - Handle irregular cases (e.g., ran → run).
3. Write down the stemmed version of the paragraph.

Paragraph-2:

"Connectivity is increasingly important in the modern digital world, connecting people and devices across continents. Businesses are streamlining their processes, enhancing collaboration, and maximizing efficiency through connected systems. These interconnected networks support applications ranging from simple email exchanges to complex data analytics, improving productivity and enabling seamless communication among teams."

Instructions

1. Tokenization: Split the paragraph into individual words (tokens).
2. Stemming: Apply the Porter Stemmer algorithm manually to each token, documenting the transformations based on the algorithm's rules. The primary steps of the Porter Stemmer Algorithm are:
 - Step 1a: Deals with plurals and simple suffixes.
 - Step 1b: Deals with suffixes like -ed or -ing.
 - Step 1c: Transforms y to i when another vowel is in the stem.

Step 2: Maps double suffices to single ones.

Step 3: Deals with suffixes like -ic, -ative, and similar.

Step 4: Removes suffixes like -al, -ance, and so on.

Step 5a: Removes a final -e.

Step 5b: Deletes -l if double -ll and preceded by a short vowel.

3. **Compose Stemmed Paragraph:** Combine the stemmed tokens back into a coherent paragraph.

Annex-A: Porter Stemmer Algorithm

Porter stemmer

The Porter stemming algorithm is a process for removing suffixes from words in English. Removing suffixes automatically is an operation which is especially useful in the field of information retrieval. In a typical IR environment a document is represented by a vector of words, or terms. Terms with a common stem will usually have similar meanings, for example:

CONNECT
CONNECTED
CONNECTING
CONNECTION
CONNECTIONS

Frequently, the performance of an IR system will be improved if term groups such as this are conflated into a single term. This may be done by removal of the various suffixes -ED, -ING, -ION, IONS to leave the single term CONNECT. In addition, the suffix stripping process will reduce the total number of terms in the IR system, and hence reduce the size and complexity of the data in the system, which is always advantageous.

Usually is desired that only inflectional morphemes are removed (those corresponding to declinations, conjugations, etc), not also derivational morphemes (which correspond to different parts of speech). Porter algorithm does not fulfill this. One can make his own set of rules (for any language). Porter's stemmer advantage is its simplicity and speed.

Porter algorithm was made in the assumption that we don't have a stem dictionary and that the purpose of the task is to improve IR performance (not as a linguistic exercise). The program is given an explicit list of suffixes, and, with each suffix, the criterion under which it may be removed from a word to leave a valid stem.

Perhaps the best criterion for removing suffixes from two words W1 and W2 to produce a single stem S, is to say that we do so if there appears to be no difference between the two statements 'a document is about W1' and 'a document is about W2'. So if W1='CONNECTION' and W2='CONNECTIONS' it seems very reasonable to conflate them to a single stem. But if W1='RELATE' and W2='RELATIVITY' it seems perhaps unreasonable, especially if the document collection is concerned with theoretical physics. Between these two extremes there is a continuum of different cases, and given two terms W1 and W2, there will be some variation in opinion as to whether they should be conflated, just as there is with deciding the relevance of some document to a query. The evaluation of the worth of a suffix stripping system is correspondingly difficult.

The success rate for the suffix stripping will be significantly less than 100% irrespective of how the process is evaluated. For example, if SAND and SANDER get conflated, so most probably will WAND and WANDER. The error here is that the -ER of WANDER has been treated as a suffix when in fact it is part of the stem. Equally, a suffix may completely alter the meaning of a word, in which case its removal is unhelpful. PROBE and PROBATE for example, have quite distinct meanings in modern English. The addition of more rules to increase the performance in one area of the vocabulary causes an equal degradation of performance elsewhere. It is also easy to give undue emphasis to cases which appear to be important, but which turn out to be rather rare. For example, cases in which the root of a word changes with the addition of a suffix, as in DECEIVE/DECEPTION, RESUME/RESUMPTION, INDEX/INDICES occur much more rarely in real vocabularies than one might at first suppose. In view of the error rate that must in any case be expected, it did not seem worthwhile to try and cope with these cases.

THE ALGORITHM

A consonant in a word is a letter other than A, E, I, O or U, and other than Y preceded by a consonant. So in TOY the consonants are T and Y, and in SYZYGY they are S, Z and G. If a letter is not a consonant it is a vowel.

A consonant will be denoted by c, a vowel by v. A list ccc... of length greater than 0 will be denoted by C, and a list vvv... of length greater than 0 will be denoted by V. Any word, or part of a word, therefore has one of the four forms:

CVCV ... C
CVCV ... V
VCVC ... C
VCVC ... V

These may all be represented by the single form:

[C]VCVC ... [V]

where the square brackets denote arbitrary presence of their contents.
Using (VC)m to denote VC repeated m times, this may again be written as:

[C] (VC){m} [V] .

m will be called the measure of any word or word part when represented in this form. The case m = 0 covers the null word. Here are some examples:

m=0 TR, EE, TREE, Y, BY.
m=1 TROUBLE, OATS, TREES, IVY.
m=2 TROUBLES, PRIVATE, OATEN, ORRERY.

The rules for removing a suffix will be given in the form:

(condition) S1 -> S2

This means that if a word ends with the suffix S1, and the stem before S1 satisfies the given condition, S1 is replaced by S2. The condition is usually given in terms of m, e.g.:

(m > 1) EMENT ->

Here S1 is 'EMENT' and S2 is null. This would map REPLACEMENT to REPLAC, since REPLAC is a word part for which m = 2.

The 'condition' part may also contain the following:

*S - the stem ends with S (and similarly for the other letters).

v - the stem contains a vowel.

*d - the stem ends with a double consonant (e.g. -TT, -SS).

*o - the stem ends cvc, where the second c is not W, X or Y (e.g. -WIL, -HOP).

And the condition part may also contain expressions with and, or and not, so that:

($m > 1$ and (*S or *T))

tests for a stem with $m > 1$ ending in S or T, while

(*d and not (*L or *S or *Z))

tests for a stem ending with a double consonant other than L, S or Z.

Elaborate conditions like this are required only rarely.

In a set of rules written beneath each other, only one is obeyed, and this will be the one with the longest matching S1 for the given word. For example, with:

SSES -> SS
IES -> I
SS -> SS
S ->

(here the conditions are all null) CARESSES maps to CARESS since SSES is the longest match for S1. Equally CARESS maps to CARESS (S1='SS') and CARES to CARE (S1='S').

In the rules below, examples of their application, successful or otherwise, are given on the right in lower case. The algorithm now follows:

Step 1a

SSES -> SS	caresses -> caress
IES -> I	ponies -> poni
	ties -> ti
SS -> SS	caress -> caress
S ->	cats -> cat

Step 1b

($m > 0$) EED -> EE	feed -> feed
	agreed -> agree
(*v*) ED ->	plastered -> plaster
	bled -> bled
(*v*) ING ->	motoring -> motor
	sing -> sing

If the second or third of the rules in Step 1b is successful, the following is done:

AT -> ATE	conflat(ed) -> conflate
BL -> BLE	troubl(ed) -> trouble
IZ -> IZE	siz(ed) -> size
(*d and not (*L or *S or *Z))	
-> single letter	
	hopp(ing) -> hop
	tann(ed) -> tan
	fall(ing) -> fall
	hiss(ing) -> hiss
	fizz(ed) -> fizz
(m=1 and *o) -> E	fail(ing) -> fail
	fil(ing) -> file

The rule to map to a single letter causes the removal of one of the double letter pair. The -E is put back on -AT, -BL and -IZ, so that the suffixes -ATE, -BLE and -IZE can be recognized later. This E may be removed in step 4.

Step 1c

(*v*) Y -> I	happy -> happi
	sky -> sky

Step 1 deals with plurals and past participles. The subsequent steps are much more straightforward.

Step 2

(m>0) ATIONAL -> ATE	relational -> relate
(m>0) TIONAL -> TION	conditional -> condition
	rational -> rational
(m>0) ENCI -> ENCE	valenci -> valence
(m>0) ANCI -> ANCE	hesitanci -> hesitance
(m>0) IZER -> IZE	digitizer -> digitize
(m>0) ABLI -> ABLE	conformabli -> conformable
(m>0) ALLI -> AL	radicalli -> radical
(m>0) ENTLI -> ENT	differentli -> different
(m>0) ELI -> E	vileli -> vile
(m>0) OUSLI -> OUS	analogousli -> analogous
(m>0) IZATION -> IZE	vietnamization -> vietnamize
(m>0) ATION -> ATE	predication -> predicate
(m>0) ATOR -> ATE	operator -> operate
(m>0) ALISM -> AL	feudalism -> feudal
(m>0) IVENESS -> IVE	decisiveness -> decisive
(m>0) FULNESS -> FUL	hopefulness -> hopeful
(m>0) OUSNESS -> OUS	callousness -> callous
(m>0) ALITI -> AL	formaliti -> formal
(m>0) IVITI -> IVE	sensitiviti -> sensitive
(m>0) BILITI -> BLE	sensibiliti -> sensible

The test for the string S1 can be made fast by doing a program switch on the penultimate letter of the word being tested. This gives a fairly even breakdown of the possible values of the string S1. It will be seen in fact that the S1-strings in step 2 are presented here in the alphabetical order of their penultimate letter.

Similar techniques may be applied in the other steps.

Step 3

(m>0) ICATE ->	IC	triplicate	->	triplic
(m>0) ATIVE ->		formative	->	form
(m>0) ALIZE ->	AL	formalize	->	formal
(m>0) ICITI ->	IC	electriciti	->	electric
(m>0) ICAL ->	IC	electrical	->	electric
(m>0) FUL ->		hopeful	->	hope
(m>0) NESS ->		goodness	->	good

Step 4

(m>1) AL ->		revival	->	reviv
(m>1) ANCE ->		allowance	->	allow
(m>1) ENCE ->		inference	->	infer
(m>1) ER ->		airliner	->	airlin
(m>1) IC ->		gyroscopic	->	gyroscop
(m>1) ABLE ->		adjustable	->	adjust
(m>1) IBLE ->		defensible	->	defens
(m>1) ANT ->		irritant	->	irrit
(m>1) EMENT ->		replacement	->	replac
(m>1) MENT ->		adjustment	->	adjust
(m>1) ENT ->		dependent	->	depend
(m>1 and (*S or *T)) ION ->		adoption	->	adopt
(m>1) OU ->		homologou	->	homolog
(m>1) ISM ->		communism	->	commun
(m>1) ATE ->		activate	->	activ
(m>1) ITI ->		angulariti	->	angular
(m>1) OUS ->		homologous	->	homolog
(m>1) IVE ->		effective	->	effect
(m>1) IZE ->		bowdlerize	->	bowdler

The suffixes are now removed. All that remains is a little tidying up.

Step 5a

(m>1) E ->		probate	->	probat
		rate	->	rate
(m=1 and not *o) E ->		cease	->	ceas

Step 5b

(m > 1 and *d and *L) ->	single letter		
	controll	->	control
	roll	->	roll

The algorithm is careful not to remove a suffix when the stem is too short, the length of the stem being given by its measure, m. There is no linguistic basis for this approach. It was merely observed that m could be used quite effectively to help decide whether or not it was wise to take off a suffix.

For example, in the following two lists:

list A	list B
-----	-----
RELATE	DERIVATE
PROBATE	ACTIVATE
CONFLATE	DEMONSTRATE
PIRATE	NECESSITATE
PRELATE	RENOVATE

-ATE is removed from the list B words, but not from the list A words. This means that the pairs DERIVATE/DERIVE, ACTIVATE/ACTIVE, DEMONSTRATE/DEMONSTRABLE, NECESSITATE/NECESSITOUS, will conflate together. The fact that no attempt is made to identify prefixes can make the results look rather inconsistent. Thus PRELATE does not lose the -ATE, but ARCHPRELATE becomes ARCHPREL.

Complex suffixes are removed bit by bit in the different steps. Thus GENERALIZATIONS is stripped to GENERALIZATION (Step 1), then to GENERALIZE (Step 2), then to GENERAL (Step 3), and then to GENER (Step 4). OSCILLATORS is stripped to OSCILLATOR (Step 1), then to OSCILLATE (Step 2), then to OSCILL (Step 4), and then to OSCIL (Step 5). In a vocabulary of 10,000 words, the reduction in size of the stem was distributed among the steps as follows:

Suffix stripping of a vocabulary of 10,000 words

Number of words reduced in step 1:	3597
" 2:	766
" 3:	327
" 4:	2424
" 5:	1373
Number of words not reduced:	3650

The resulting vocabulary of stems contained 6370 distinct entries. Thus the suffix stripping process reduced the size of the vocabulary by about one third.

Similar algorithms can be applied for different languages, one can define his own rules adapted to his own needs.